

Data Carpentry

The Craft of Working with Data

Brendan Knapp and Christopher Callaghan

2020-10-03

Contents

| | |
|---|----------|
| Welcome | 3 |
| Preface | 4 |
| I Setup | 5 |
| 1 R and RStudio | 6 |
| 1.1 R | 6 |
| 1.1.1 Installation | 6 |
| 1.2 RStudio | 6 |
| 1.2.1 Installation | 6 |
| 2 The Basics | 7 |
| 2.1 R as a Calculator | 7 |
| 2.2 Fundamental Types | 8 |
| 2.3 Variables | 9 |
| 2.4 Multiple Values | 10 |
| 2.5 Functions | 10 |
| 2.6 Documentation | 12 |
| 2.7 Missingness and Nothingness | 13 |
| 2.7.1 NA | 13 |
| 2.7.2 NULL | 13 |
| 2.8 Predicate Functions | 13 |
| 2.9 Vectorized Functions | 15 |

| | | |
|---------|----------------------------------|----|
| II | Reading and Writing Data | 17 |
| 3 | Tabular Data | 18 |
| 3.1 | Basics | 18 |
| 3.2 | Common Pitfalls | 22 |
| 3.2.1 | Incorrect Column Types | 22 |
| 3.2.1.1 | Solution | 22 |
| III | Data Frames | 24 |
| 4 | Manipulating Data Frames | 25 |
| 4.1 | select() Columns | 26 |
| 4.1.1 | by Name | 26 |
| 4.1.2 | by Index | 27 |
| 4.1.3 | by Name Pattern | 29 |
| 4.1.4 | by Data Type | 30 |
| 4.2 | filter() Rows | 32 |
| 4.2.1 | by row_number() | 32 |
| 4.2.2 | by Name | 33 |
| 4.2.3 | by Type | 34 |
| 4.3 | arrange() Rows | 35 |

Welcome

Test

<- == !=

```
test <- "face"
```

Preface

init

Part I

Setup

Chapter 1

R and RStudio

1.1 R

1.1.1 Installation

<https://cran.r-project.org/>

1.2 RStudio

1.2.1 Installation

<https://rstudio.com/products/rstudio/download/>

Chapter 2

The Basics

```
"Hello, World!"  
#> [1] "Hello, World!"
```

2.1 R as a Calculator

```
1 + 1                # addition  
#> [1] 2  
  
1 - 1                # subtraction  
#> [1] 0  
  
2 * 3                # multiplication  
#> [1] 6  
  
2 / 3                # division  
#> [1] 0.667  
  
1 + 1 * 3            # combining operations  
#> [1] 4  
  
(1 + 1) * 3          # operator precedence  
#> [1] 6  
  
3 / 2                # division  
#> [1] 1.5  
  
#  $\downarrow$  pronounced "modulo"  
3 %% 2               # division remainder  
#> [1] 1
```



```

4 %/% 2          # integer division
#> [1] 2

3^2             # exponents
#> [1] 9
4**2           # also exponents!
#> [1] 16

Inf + 1         #
#> [1] Inf

```



If your code doesn't form a complete *expression*, then R will look for the rest of on the next line. Here's an example:

```
1 +
```

1 + isn't a complete expression, so R will look for more on subsequent lines. You'll see something like the following:

```

> 1 +
+
+
+

```

If this happens, press the **Esc**(scape) key (you may have to click on the Console pane first) and fix your code.

2.2 Fundamental Types

R has several basic data types that serve as the foundation upon which everything is built.

```

1          # double (short for double-precision floating-point number)
#> [1] 1
3.14      # also double (we can just think of them as decimals)
#> [1] 3.14

1L        # integer (`L` for "Literal" or `long` integers)
#> [1] 1

"1"       # character or string (kinda... we'll discuss later)
#> [1] "1"

TRUE      # logical (similar to `bool`s in other languages)
#> [1] TRUE
FALSE     # also logical
#> [1] FALSE

```

```
4i          # complex (we're never going to use these)
#> [1] 0+4i
```



Like most programming languages, R lets us mix *comments* into our code. Anything that follows `#` on the same line is ignored by R.

Comments enable us to annotate our work or temporarily (hopefully) disable lines of code.

```
-1 * -1000 # a negative number times a negative is positive
#> [1] 1000

# TRUE + FALSE # felt cute, might un-comment later
```

Leverage comments to communicate with humans! They're an opportunity for explaining *what* something does and (often more-importantly) *why* something works or is necessary.

Since comments are ubiquitous, it's worth pointing out two common conventions:

```
# TODO(CC): CC (initials) is going to implement some behavior
# FIXME(BK): BK is going to fix some problem
```

2.3 Variables

R's assignment operator is `<-`.

```
my_first_var <- "referring to data w/ names is handy!"
my_first_var
#> [1] "referring to data w/ names is handy!"
```

We can also use `=` like many other languages, but we *highly* discourage this (especially starting out) because we use `=` elsewhere. If you stick to `<-`, you'll never have to guess where you've assigned variables or rely on context clues to predict `=`'s intended purpose or behavior.

You should always prefer descriptive variable names so that others can more easily understand your code. Most of the time, the other person will just be you in the future.

```
length <- 2
width <- 4

area <- length * width
area
#> [1] 8
```

2.4 Multiple Values

We'll almost always need to deal with more than one value, so R let's us `c()`ombine values.

```
c(1, 2, 3, 4, 5)
#> [1] 1 2 3 4 5
```

We won't get into the nitty-gritty details just yet, but we typically call a collection of values of the same type (*homogeneous*) a **vector**.

R is special for a few reasons and having native **vectors** is definitely one of them. Understanding how they work is fundamental to writing good (and fast!) code.

```
c(1, 2, 3, 4, 5, 6)      # `c()` is short for "combine"
#> [1] 1 2 3 4 5 6

1:6                      # `:` lets us create sequences
#> [1] 1 2 3 4 5 6

my_first_vector <- -5:5 # we'll explain `vector`s later,
my_first_vector
#> [1] -5 -4 -3 -2 -1 0 1 2 3 4 5
```

2.5 Functions

```
# ↓ name of function
sqrt(x = 16)
#> [1] 4
```

R comes with some handy variables built-in , such as `letters` and `LETTERS`.

```
letters
#> [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "t" "u" "v"
#> [23] "w" "x" "y" "z"

LETTERS
#> [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S" "T" "U" "V"
#> [23] "W" "X" "Y" "Z"
```

```
#      ↓ parameter, formal, or argument
toupper(x = letters)
#> [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S" "T" "U" "V"
#> [23] "W" "X" "Y" "Z"
```

```
#          ↓↓↓↓↓↓ argument (always)
tolower(x = LETTERS)
#> [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "t" "u" "v"
#> [23] "w" "x" "y" "z"
```

We refer to `x = letters` as a *named* argument because we specify the parameter (`x`) to which we're passing our argument (`letters`), but we often don't specify the name of a parameter.

```
tolower(letters)
#> [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s" "t" "u" "v"
#> [23] "w" "x" "y" "z"
```

We can't screw up too easily since `tolower()` and `toupper()` only have one parameter (`x`), but many functions can take multiple arguments.

Let's say we have a vector of `unsorted_numbers`:

```
unsorted_numbers <- c(3, 2, 10, 8, 1, 4, 9, 6, 5, 7)
unsorted_numbers
#> [1] 3 2 10 8 1 4 9 6 5 7
```

Like most languages, R has a built-in `sort()` function we can use, which works like so:

```
sort(x = unsorted_numbers)
#> [1] 1 2 3 4 5 6 7 8 9 10
```

By default, `sort()` sorts in *ascending* order, but we oftentimes will want to sort in *descending* (or *decreasing*) order.

Rather than having a separate function called `sort_decreasing()`, we pass an argument to `sort()`'s `decreasing` parameter.

```
sort(x = unsorted_numbers, decreasing = TRUE)
#> [1] 10 9 8 7 6 5 4 3 2 1
```

Even though `sort()` has multiple parameters, we can still skip the names if we pass our arguments *by position*.

```
sort(unsorted_numbers, TRUE)
#> [1] 10 9 8 7 6 5 4 3 2 1
```

Considering that `x` is `sort()`'s first parameter, and `decreasing` is `sort()`'s second parameter, we can pass our arguments (`unsorted_numbers` and `TRUE`) in the same order and R will know what we meant.

We can also mix *positional* and *named* argument (and often do), but we should always prioritize *readable* code.

```

sort(unsorted_numbers, decreasing = TRUE) # good
#> [1] 10 9 8 7 6 5 4 3 2 1

sort(x = unsorted_numbers, TRUE)          # avoid this
#> [1] 10 9 8 7 6 5 4 3 2 1

sort(decreasing = TRUE, unsorted_numbers) # just... no
#> [1] 10 9 8 7 6 5 4 3 2 1

```

2.6 Documentation

You're hopefully wondering "How could we know the order of `sort()`'s parameters?" which leads us to documentation. If you want more information on a specific function, you should check out the documentation, which you can do with `?` or `help()`. Here's what that looks like for `sort()`

```
?sort
```

```
help(sort) # does the same thing as `?sort`
```

```

sort {base}
R Documentation

Sorting or Ordering Vectors

Description
Sort (or order) a vector or factor (partially) into ascending or descending order. For ordering along more than one variable, e.g., for sorting data frames, see
order.

Usage
sort(x, decreasing = FALSE, ...)

## Default S3 method:
sort(x, decreasing = FALSE, na.last = NA, ...)

sort.int(x, partial = NULL, na.last = NA, decreasing = FALSE,
method = c("auto", "shell", "quick", "radix"), index.return = FALSE)

Arguments
x
    for sort an R object with a class or a numeric, complex, character or logical vector. For sort.int, a numeric, complex, character or logical
    vector, or a factor.

decreasing
    logical. Should the sort be increasing or decreasing? For the "radix" method, this can be a vector of length equal to the number of arguments
    in .... For the other methods, it must be length one. Not available for partial sorting.

...
    arguments to be passed to or from methods or (for the default methods and objects without a class) to sort.int.

na.last
    for controlling the treatment of NAs. If TRUE, missing values in the data are put last; if FALSE, they are put first; if NA, they are removed.

partial
    NULL or a vector of indices for partial sorting.

method
    character string specifying the algorithm used. Not available for partial sorting. Can be abbreviated.

index.return
    logical indicating if the ordering index vector should be returned as well. Supported by method == "radix" for any na.last mode and data
    type, and the other methods when na.last = NA (the default) and fully sorting non-factors.

```

There's *a ton* of information here, but all we're interested in at the moment is the order in which we need to pass arguments to `sort()`, which we can find in the **Arguments** section.

2.7 Missingness and Nothingness

2.7.1 NA

You may have noticed the `na.last` argument in `sort()`'s documentation. R can represent nothingness with `NULL` (or `None` in Python), but it can also represent *unknown* or *missing* values with `NA`.

```
unsorted_numbers_with_nas <- c(3, 2, 10, 8, NA, 1, 4, 9, NA, 6, 5, 7)
unsorted_numbers_with_nas
#> [1] 3 2 10 8 NA 1 4 9 NA 6 5 7
```

`sort()`'s default behavior is `na.last = NA`, which simply removes any NAs.

```
sort(unsorted_numbers_with_nas)
#> [1] 1 2 3 4 5 6 7 8 9 10
```

If we want to keep NAs, we must specify whether `sort()` places them first or last.

```
sort(unsorted_numbers_with_nas, na.last = TRUE)
#> [1] 1 2 3 4 5 6 7 8 9 10 NA NA
sort(unsorted_numbers_with_nas, na.last = FALSE)
#> [1] NA NA 1 2 3 4 5 6 7 8 9 10
```

2.7.2 NULL

For most purposes, the difference between `NA` and `NULL` is that vectors (like `unsorted_numbers_with_nas`) can have `NA` values.

If we try to put `NULL` in a vector, it simply disappears.

```
c(3, 2, 10, 8, NULL, 1, 4, 9, NULL, 6, 5, 7)
#> [1] 3 2 10 8 1 4 9 6 5 7
```

But, how do we check if something is `NA` or `NULL`?

2.8 Predicate Functions

A *predicate function* is a function that returns either `TRUE` or `FALSE` based on some condition the function is checking.

Predicate function (should) use a name that expresses this intent, such as `is<some condition>`, `any<some condition>()`, or `all<some condition>()`.

If we want check if something is `NULL`, we use `is.null()`.

```
is.null("this string isn't NULL!")
#> [1] FALSE
is.null(NULL)
#> [1] TRUE
```

R has *many* built-in predicate functions, including ones to check the basic data types that we've already seen.

```
is.double(1)
#> [1] TRUE
is.double(1L)
#> [1] FALSE

vec_dbl <- c(8, 6, 7, 5, 3, 0, 9)
is.double(vec_dbl)
#> [1] TRUE

is.integer(1)
#> [1] FALSE
is.integer(1L)
#> [1] TRUE
vec_int <- 1:10
is.integer(vec_int)
#> [1] TRUE

is.character(3.14)
#> [1] FALSE
is.character("is it though?")
#> [1] TRUE
is.character(letters)
#> [1] TRUE

is.logical("the year 2020")
#> [1] FALSE
is.logical(TRUE)
#> [1] TRUE
is.logical(FALSE)
#> [1] TRUE
vec_lgl <- c(TRUE, FALSE, TRUE)
is.logical(vec_lgl)
#> [1] TRUE
```

Similar to `is.null()`, there's `is.na()`.

```
is.na("not NA!")
#> [1] FALSE
is.na(NA)
#> [1] TRUE
```

Recall our variable `unsorted_numbers_with_nas`.

```
unsorted_numbers_with_nas
#> [1] 3 2 10 8 NA 1 4 9 NA 6 5 7
```

Consider the following:

- The predicate functions we’ve seen so far **return** either **TRUE** or **FALSE**.
- **vectors** can contain both **NA** *and* non-NA values.

Can you guess what `is.na()` returns?

```
is.na(unsorted_numbers_with_nas)
#> [1] FALSE FALSE FALSE FALSE TRUE FALSE FALSE FALSE TRUE FALSE FALSE FALSE
```



We’ll discuss accessing a **vector**’s individual elements later, but `is.na()` is what we call a *vectorized* function: a function that takes **vector** argument and operates every on element simultaneously.

2.9 Vectorized Functions

As high-speed R coders, we should prefer **vectorized** solutions whenever possible as they’re not only idiomatic (and thus easy for other R users to understand), but they’re typically several orders of magnitude faster than other solutions.

While R isn’t the fastest language out there, complaints about its speed often come poor code, which is often “speaking” R with a C or Python accent.

The simplest way to wrap our heads around **vectorized** operations probably is with math.

To do this, let’s first make a **vector** with five 0s in it.

We could do that like the following:


```
c(0, 0, 0, 0, 0)
#> [1] 0 0 0 0 0
```

But, good coders are lazy coders who want to (correctly) automate everything they can. With that in mind, let's `rep()`eat 0 5 times.

```
zeros <- rep(0, length = 5)
zeros
#> [1] 0 0 0 0 0
```

For our purposes, the term *scalar* refers to an object that is a single value.

If we want to add 1 (a scalar) to every element of `zeros`, we can run `zeros + 1` or `1 + zeros`:

```
zeros + 1
#> [1] 1 1 1 1 1
```

R knows that 1 is a single value (and assumes we know what we're doing) and performs the operation (+) between it and every element of `zeros`. In R-speak, we refer to this behavior as *recycling*.

Let's see what happens when we add `zeros` and a *vector* containing two elements.

```
two_threes <- c(3, 3)

zeros + two_threes
#> Warning in zeros + two_threes: longer object length is not a multiple of shorter object length
#> [1] 3 3 3 3 3
```

That's probably not what you expected, but R threw a `warning()` to tell us something seems wrong.

R let's us get away with a lot of things it shouldn't, which includes

Part II

Reading and Writing Data

Chapter 3

Tabular Data

- Aliases:
 - Tabular files
 - Flat
 - Delimited
- Includes:
 - Comma-Separated Value (.csv)
 - Tab-Separated Value (.tsv)

3.1 Basics

```
library(readr)
```

Here's some example data, modified from <http://www.gapminder.org/data/>

```
country,continent,year,lifeExp,pop,gdpPercap      # header/column names, separated by commas
Afghanistan,Asia,1952,28.801,8425333,779.4453145
Afghanistan,Asia,1957,30.332,9240934,820.8530296  # comma-separated values
Afghanistan,Asia,1962,31.997,10267083,853.10071
Afghanistan,Asia,1967,34.02,11537966,836.1971382
Afghanistan,Asia,1972,36.088,13079460,739.9811058
Afghanistan,Asia,1977,38.438,14880372,786.11336
Afghanistan,Asia,1982,39.854,12881816,978.0114388
Afghanistan,Asia,1987,40.822,13867957,852.3959448
```

```
csv_text <-
'country,continent,year,lifeExp,pop,gdpPercap
Afghanistan,Asia,1952,28.801,8425333,779.4453145
Afghanistan,Asia,1957,30.332,9240934,820.8530296
Afghanistan,Asia,1962,31.997,10267083,853.10071'
```

```
Afghanistan,Asia,1967,34.02,11537966,836.1971382
Afghanistan,Asia,1972,36.088,13079460,739.9811058
Afghanistan,Asia,1977,38.438,14880372,786.11336
Afghanistan,Asia,1982,39.854,12881816,978.0114388
Afghanistan,Asia,1987,40.822,13867957,852.3959448'
```

```
csv_file <- tempfile(fileext = ".csv")
csv_file # a temporary file path
#> [1] "/tmp/Rtmp6lDJkS/file73aac04c4a3.csv"
writeLines(text = csv_text, con = csv_file) # write `csv_text` to `csv_file`
```

```
read_csv(file = csv_file)
#> Parsed with column specification:
#> cols(
#>   country = col_character(),
#>   continent = col_character(),
#>   year = col_double(),
#>   lifeExp = col_double(),
#>   pop = col_double(),
#>   gdpPercap = col_double()
#> )
#> # A tibble: 8 x 6
#>   country      continent  year lifeExp      pop gdpPercap
#>   <chr>         <chr>    <dbl>  <dbl>    <dbl>    <dbl>
#> 1 Afghanistan Asia      1952   28.8  8425333    779.
#> 2 Afghanistan Asia      1957   30.3  9240934    821.
#> 3 Afghanistan Asia      1962   32.0 10267083    853.
#> 4 Afghanistan Asia      1967   34.0 11537966    836.
#> 5 Afghanistan Asia      1972   36.1 13079460    740.
#> 6 Afghanistan Asia      1977   38.4 14880372    786.
#> 7 Afghanistan Asia      1982   39.9 12881816    978.
#> 8 Afghanistan Asia      1987   40.8 13867957    852.
```

You may encounter Tab-Delimited data where values are separated by `\t` instead of `,`. Instead of `readr::read_csv()`, we can use `readr::read_tsv()`.

```
tsv_text <-
'country\tcontinent\tyear\tlifeExp\tpop\tgdpPercap
Afghanistan\tAsia\t1952\t28.801\t8425333\t779.4453145
Afghanistan\tAsia\t1957\t30.332\t9240934\t820.8530296
Afghanistan\tAsia\t1962\t31.997\t10267083\t853.10071
Afghanistan\tAsia\t1967\t34.02\t11537966\t836.1971382
Afghanistan\tAsia\t1972\t36.088\t13079460\t739.9811058
Afghanistan\tAsia\t1977\t38.438\t14880372\t786.11336
Afghanistan\tAsia\t1982\t39.854\t12881816\t978.0114388
Afghanistan\tAsia\t1987\t40.822\t13867957\t852.3959448'

tsv_file <- tempfile(fileext = ".tsv")
writeLines(text = tsv_text, con = tsv_file)
```

```

read_tsv(file = tsv_file)
#> Parsed with column specification:
#> cols(
#>   country = col_character(),
#>   continent = col_character(),
#>   year = col_double(),
#>   lifeExp = col_double(),
#>   pop = col_double(),
#>   gdpPercap = col_double()
#> )
#> # A tibble: 8 x 6
#>   country    continent  year lifeExp      pop gdpPercap
#>   <chr>      <chr>    <dbl>  <dbl>    <dbl>    <dbl>
#> 1 Afghanistan Asia      1952   28.8  8425333    779.
#> 2 Afghanistan Asia      1957   30.3  9240934    821.
#> 3 Afghanistan Asia      1962   32.0 10267083    853.
#> 4 Afghanistan Asia      1967   34.0 11537966    836.
#> 5 Afghanistan Asia      1972   36.1 13079460    740.
#> 6 Afghanistan Asia      1977   38.4 14880372    786.
#> 7 Afghanistan Asia      1982   39.9 12881816    978.
#> 8 Afghanistan Asia      1987   40.8 13867957    852.

```

If we find ourselves reading delimited data that uses something other than `\t` or `,` to separate values, we can use `readr::read_delim()`.

```

pipe_separated_values_text <-
'country|continent|year|lifeExp|pop|gdpPercap
Afghanistan|Asia|1952|28.801|8425333|779.4453145
Afghanistan|Asia|1957|30.332|9240934|820.8530296
Afghanistan|Asia|1962|31.997|10267083|853.10071
Afghanistan|Asia|1967|34.02|11537966|836.1971382
Afghanistan|Asia|1972|36.088|13079460|739.9811058
Afghanistan|Asia|1977|38.438|14880372|786.11336
Afghanistan|Asia|1982|39.854|12881816|978.0114388
Afghanistan|Asia|1987|40.822|13867957|852.3959448'

psv_file <- tempfile(fileext = ".tsv")
writeLines(text = pipe_separated_values_text, con = psv_file)

```

```

read_delim(file = psv_file, delim = "|")
#> Parsed with column specification:
#> cols(
#>   country = col_character(),
#>   continent = col_character(),
#>   year = col_double(),
#>   lifeExp = col_double(),

```

```
#> pop = col_double(),
#> `gdpPercap` = col_double()
#> )
#> # A tibble: 8 x 6
#>   country      continent year lifeExp      pop `gdpPercap`
#>   <chr>         <chr>    <dbl>  <dbl>    <dbl>      <dbl>
#> 1 Afghanistan Asia      1952   28.8  8425333    779.
#> 2 Afghanistan Asia      1957   30.3  9240934    821.
#> 3 Afghanistan Asia      1962   32.0 10267083    853.
#> 4 Afghanistan Asia      1967   34.0 11537966    836.
#> 5 Afghanistan Asia      1972   36.1 13079460    740.
#> 6 Afghanistan Asia      1977   38.4 14880372    786.
#> 7 Afghanistan Asia      1982   39.9 12881816    978.
#> 8 Afghanistan Asia      1987   40.8 13867957    852.
```

```
country,continent,year,lifeExp,pop,gdpPercap      # header/column names
Afghanistan,Asia,1952,28.801,8425333,779.4453145
Afghanistan,Asia,1957,30.332,9240934,820.8530296
Afghanistan,Asia,1962,31.997,10267083,853.10071
Afghanistan,Asia,1967,34.02,11537966,836.1971382
Afghanistan,Asia,1972,36.088,13079460,739.9811058
Afghanistan,Asia,1977,38.438,14880372,786.11336
Afghanistan,Asia,1982,39.854,12881816,978.0114388
Afghanistan,Asia,1987,40.822,13867957,852.3959448
Afghanistan,,N/A,,                                # notice that we're missing values
```

```
csv_text <-
'country,continent,year,lifeExp,pop,gdpPercap
Afghanistan,Asia,1952,28.801,8425333,779.4453145
Afghanistan,Asia,1957,30.332,9240934,820.8530296
Afghanistan,Asia,1962,31.997,10267083,853.10071
Afghanistan,Asia,1967,34.02,11537966,836.1971382
Afghanistan,Asia,1972,36.088,13079460,739.9811058
Afghanistan,Asia,1977,38.438,14880372,786.11336
Afghanistan,Asia,1982,39.854,12881816,978.0114388
Afghanistan,Asia,1987,40.822,13867957,852.3959448
Afghanistan,,N/A,, '

csv_file <- tempfile(fileext = ".csv")
writeLines(text = csv_text, con = csv_file)
```

3.2 Common Pitfalls

3.2.1 Incorrect Column Types

```
data_frame_from_csv <- read_csv(file = csv_file)
#> Parsed with column specification:
#> cols(
#>   country = col_character(),
#>   continent = col_character(),
#>   year = col_double(),
#>   lifeExp = col_character(),
#>   pop = col_double(),
#>   gdpPercap = col_double()
#> )
data_frame_from_csv
#> # A tibble: 9 x 6
#>   country      continent  year lifeExp      pop gdpPercap
#>   <chr>         <chr>    <dbl> <chr>      <dbl>    <dbl>
#> 1 Afghanistan Asia      1952 28.801  8425333  779.
#> 2 Afghanistan Asia      1957 30.332  9240934  821.
#> 3 Afghanistan Asia      1962 31.997 10267083  853.
#> 4 Afghanistan Asia      1967 34.02  11537966  836.
#> 5 Afghanistan Asia      1972 36.088 13079460  740.
#> 6 Afghanistan Asia      1977 38.438 14880372  786.
#> 7 Afghanistan Asia      1982 39.854 12881816  978.
#> 8 Afghanistan Asia      1987 40.822 13867957  852.
#> 9 Afghanistan <NA>      NA N/A      NA      NA
```

Notice that our `year` column says `<dbl>`, referring to it being of type `double`, yet all of our `year` values are whole numbers.

```
typeof(data_frame_from_csv$year)
#> [1] "double"
data_frame_from_csv$year
#> [1] 1952 1957 1962 1967 1972 1977 1982 1987 NA
```

We also have "N/A" in our `lifeExp` column, forcing R to interpret all `lifeExp` values as `characters` (`<chr>`).

```
typeof(data_frame_from_csv$lifeExp)
#> [1] "character"
data_frame_from_csv$lifeExp
#> [1] "28.801" "30.332" "31.997" "34.02" "36.088" "38.438" "39.854" "40.822" "N/A"
```

3.2.1.1 Solution

```

read_csv(
  file = csv_file,
  col_types = cols(
    country = col_character(),
    continent = col_character(),
    year = col_integer(),      # read `year` as `integer`
    lifeExp = col_double(),    # read `lifeExp` as `double`
    pop = col_double(),
    gdpPercap = col_double()
  ),
  na = c("", "N/A")          # be explicit about how `csv_file` represents missing values
)

#> # A tibble: 9 x 6
#>   country    continent  year lifeExp      pop gdpPercap
#>   <chr>      <chr>    <int>  <dbl>    <dbl>    <dbl>
#> 1 Afghanistan Asia      1952   28.8  8425333    779.
#> 2 Afghanistan Asia      1957   30.3  9240934    821.
#> 3 Afghanistan Asia      1962   32.0 10267083    853.
#> 4 Afghanistan Asia      1967   34.0 11537966    836.
#> 5 Afghanistan Asia      1972   36.1 13079460    740.
#> 6 Afghanistan Asia      1977   38.4 14880372    786.
#> 7 Afghanistan Asia      1982   39.9 12881816    978.
#> 8 Afghanistan Asia      1987   40.8 13867957    852.
#> 9 Afghanistan <NA>      NA      NA      NA      NA

```


Part III

Data Frames

Chapter 4

Manipulating Data Frames

```
library(tidyverse, warn.conflicts = FALSE)
#> -- Attaching packages ----- tidyverse 1.3.0 --
#> v ggplot2 3.3.2      v purrr  0.3.4
#> v tibble  3.0.3      v dplyr  1.0.2
#> v tidyr   1.1.2      v stringr 1.4.0
#> v readr   1.3.1      v forcats 0.5.0
#> -- Conflicts ----- tidyverse_conflicts() --
#> x dplyr::filter() masks stats::filter()
#> x dplyr::lag()     masks stats::lag()

df <- tibble(
  group = c("a", "a", "b", "b", "b"),
  a = c(1, 4, NA, 3, 5),
  b = c(9, NA, 8, 10, 7),
  c = c(TRUE, FALSE, NA, FALSE, TRUE),
  d = c(LETTERS[1:3], NA, LETTERS[[5]]),
  e = factor(1:5, labels = c("tiny", "small", "medium", "big", "huge")),
  f_col = c(as.Date(NA), as.Date("2020-09-23") + c(3, 2, 1, 4)),
  g_col = c(as.POSIXct("2020-09-23 00:00:00") + 1:4 * 60 * 60 * 24 * 1.1, NA),
  col_h = list(c(1, 10), c(2, NA), c(3, 8), c(4, 7), c(5, 6)),
  col_i = list(NULL, pi, month.abb[6:10], iris, as.matrix(mtcars))
)

df
#> # A tibble: 5 x 10
#>   group     a     b     c     d     e     f_col       g_col       col_h     col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>       <dtm>       <list> <list>
#> 1 a         1     9 TRUE   A     tiny   NA         2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 2 a         4    NA FALSE B     small  2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 3 b        NA     8 NA     C     medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
#> 4 b         3    10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 5 b         5     7 TRUE   E     huge   2020-09-27 NA         <dbl [2~ <dbl[,11] [32 x~
glimpse(df)
#> Rows: 5
```

```
#> Columns: 10
#> $ group <chr> "a", "a", "b", "b", "b"
#> $ a      <dbl> 1, 4, NA, 3, 5
#> $ b      <dbl> 9, NA, 8, 10, 7
#> $ c      <lgl> TRUE, FALSE, NA, FALSE, TRUE
#> $ d      <chr> "A", "B", "C", NA, "E"
#> $ e      <fct> tiny, small, medium, big, huge
#> $ f_col  <date> NA, 2020-09-26, 2020-09-25, 2020-09-24, 2020-09-27
#> $ g_col  <dtm> 2020-09-24 02:24:00, 2020-09-25 04:48:00, 2020-09-26 07:12:00, 2020-09-27 0...
#> $ col_h  <list> [<1, 10>, <2, NA>, <3, 8>, <4, 7>, <5, 6>]
#> $ col_i  <list> [NULL, 3.14, <"Jun", "Jul", "Aug", "Sep", "Oct">, <data.frame[150 x 5]>, <m...
```

4.1 select() Columns

4.1.1 by Name

```
df %>%
  select(a)
#> # A tibble: 5 x 1
#>       a
#>   <dbl>
#> 1     1
#> 2     4
#> 3    NA
#> 4     3
#> 5     5
```

```
df %>%
  select(a, c, e)
#> # A tibble: 5 x 3
#>       a c     e
#>   <dbl> <lgl> <fct>
#> 1     1 TRUE tiny
#> 2     4 FALSE small
#> 3    NA NA   medium
#> 4     3 FALSE big
#> 5     5 TRUE huge
```

```
df %>%
  select(b, d, f_col)
#> # A tibble: 5 x 3
#>       b d     f_col
#>   <dbl> <chr> <date>
#> 1     9 A     NA
```

```
#> 2    NA B    2020-09-26
#> 3     8 C    2020-09-25
#> 4    10 <NA> 2020-09-24
#> 5     7 E    2020-09-27
```

```
df %>%
  select(b, c, everything())
#> # A tibble: 5 x 10
#>       b c    group   a d     e   f_col   g_col           col_h   col_i
#>   <dbl> <lgl> <chr> <dbl> <chr> <fct> <date>   <dtm>           <list> <list>
#> 1     9 TRUE   a       1 A    tiny   NA       2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 2    NA FALSE a       4 B    small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 3     8 NA     b      NA C    medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
#> 4    10 FALSE b       3 <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 5     7 TRUE   b       5 E    huge    2020-09-27 NA           <dbl [2~ <dbl[,11] [32 x ~
```

```
df %>%
  select(b, c, everything(), -a)
#> # A tibble: 5 x 9
#>       b c    group d     e   f_col   g_col           col_h   col_i
#>   <dbl> <lgl> <chr> <chr> <fct> <date>   <dtm>           <list> <list>
#> 1     9 TRUE   a     A    tiny   NA       2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2    NA FALSE a     B    small 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3     8 NA     b     C    medium 2020-09-25 2020-09-26 07:12:00 <dbl [2]> <chr [5]>
#> 4    10 FALSE b    <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x 5]>
#> 5     7 TRUE   b     E    huge    2020-09-27 NA           <dbl [2]> <dbl[,11] [32 x 11]>
```

```
cols_to_select <- c("a", "c", "e")
df %>%
  select(all_of(cols_to_select))
#> # A tibble: 5 x 3
#>       a c     e
#>   <dbl> <lgl> <fct>
#> 1     1 TRUE tiny
#> 2     4 FALSE small
#> 3    NA NA    medium
#> 4     3 FALSE big
#> 5     5 TRUE huge
```

4.1.2 by Index

```
df %>%
  select(1L)
#> # A tibble: 5 x 1
```

```
#> group
#> <chr>
#> 1 a
#> 2 a
#> 3 b
#> 4 b
#> 5 b
```

```
df %>%
  select(1, 3, 5)
#> # A tibble: 5 x 3
#>   group      b d
#>   <chr> <dbl> <chr>
#> 1 a         9 A
#> 2 a        NA B
#> 3 b         8 C
#> 4 b        10 <NA>
#> 5 b         7 E
```

```
df %>%
  select(2, 4, 6)
#> # A tibble: 5 x 3
#>       a c      e
#>   <dbl> <lgl> <fct>
#> 1     1 TRUE tiny
#> 2     4 FALSE small
#> 3    NA NA  medium
#> 4     3 FALSE big
#> 5     5 TRUE  huge
```

```
df %>%
  select(2:3, everything())
#> # A tibble: 5 x 10
#>       a      b group c      d      e      f_col      g_col      col_h      col_i
#>   <dbl> <dbl> <chr> <lgl> <chr> <fct> <date>      <dtm>      <list> <list>
#> 1     1     9 a    TRUE A    tiny NA      2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 2     4    NA a    FALSE B    small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 3    NA     8 b    NA   C    medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
#> 4     3    10 b    FALSE <NA> big 2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 5     5     7 b    TRUE  E    huge 2020-09-27 NA      <dbl [2~ <dbl[,11] [32 x~
```

```
df %>%
  select(2:3, everything(), -1)
#> # A tibble: 5 x 9
#>       a      b c      d      e      f_col      g_col      col_h      col_i
#>   <dbl> <dbl> <lgl> <chr> <fct> <date>      <dtm>      <list> <list>
```

```
#> 1      1      9 TRUE A      tiny NA      2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2      4      NA FALSE B      small 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3      NA      8 NA      C      medium 2020-09-25 2020-09-26 07:12:00 <dbl [2]> <chr [5]>
#> 4      3      10 FALSE <NA> big 2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x 5]>
#> 5      5      7 TRUE E      huge 2020-09-27 NA <dbl [2]> <dbl[,11] [32 x 11]>
```

```
cols_to_select <- c(1, 3, 5)
df %>%
  select(all_of(cols_to_select))
#> # A tibble: 5 x 3
#>   group      b d
#>   <chr> <dbl> <chr>
#> 1 a      9 A
#> 2 a      NA B
#> 3 b      8 C
#> 4 b     10 <NA>
#> 5 b      7 E
```

```
cols_to_select <- c(1, 3, 5, 1000)
df %>%
  select(any_of(cols_to_select))
#> # A tibble: 5 x 3
#>   group      b d
#>   <chr> <dbl> <chr>
#> 1 a      9 A
#> 2 a      NA B
#> 3 b      8 C
#> 4 b     10 <NA>
#> 5 b      7 E
```

4.1.3 by Name Pattern

`contains()` selects a column if *any* part of its name contains `match=`.

```
df %>%
  select(contains(match = "col"))
#> # A tibble: 5 x 4
#>   f_col      g_col      col_h      col_i
#>   <date>    <dtm>    <list>    <list>
#> 1 NA      2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3 2020-09-25 2020-09-26 07:12:00 <dbl [2]> <chr [5]>
#> 4 2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x 5]>
#> 5 2020-09-27 NA <dbl [2]> <dbl[,11] [32 x 11]>
```

`starts_with()` selects a column if its name starts with `match=`.

```
df %>%
  select(starts_with("col_"))
#> # A tibble: 5 x 2
#>   col_h      col_i
#>   <list>    <list>
#> 1 <dbl [2]> <NULL>
#> 2 <dbl [2]> <dbl [1]>
#> 3 <dbl [2]> <chr [5]>
#> 4 <dbl [2]> <df[,5] [150 x 5]>
#> 5 <dbl [2]> <dbl[,11] [32 x 11]>
```

`starts_with()` selects a column if its name ends with `match=`.

```
df %>%
  select(ends_with("_col"))
#> # A tibble: 5 x 2
#>   f_col      g_col
#>   <date>    <dtm>
#> 1 NA       2020-09-24 02:24:00
#> 2 2020-09-26 2020-09-25 04:48:00
#> 3 2020-09-25 2020-09-26 07:12:00
#> 4 2020-09-24 2020-09-27 09:36:00
#> 5 2020-09-27 NA
```

`matches()`s Selects a column if its name matches a regular expression pattern.

```
df %>%
  select(matches("(^\\w_)?col(_\\w)?"))
#> # A tibble: 5 x 4
#>   f_col      g_col      col_h      col_i
#>   <date>    <dtm>    <list>    <list>
#> 1 NA       2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3 2020-09-25 2020-09-26 07:12:00 <dbl [2]> <chr [5]>
#> 4 2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x 5]>
#> 5 2020-09-27 NA             <dbl [2]> <dbl[,11] [32 x 11]>
```

4.1.4 by Data Type

```
df %>%
  select(where(is.factor))
#> # A tibble: 5 x 1
#>   e
#>   <fct>
#> 1 tiny
```

```
#> 2 small
#> 3 medium
#> 4 big
#> 5 huge
```

```
df %>%
  select_if(is.factor)
#> # A tibble: 5 x 1
#>   e
#>   <fct>
#> 1 tiny
#> 2 small
#> 3 medium
#> 4 big
#> 5 huge
```

```
df %>%
  select(where(is.factor), f_col)
#> # A tibble: 5 x 2
#>   e      f_col
#>   <fct>   <date>
#> 1 tiny   NA
#> 2 small 2020-09-26
#> 3 medium 2020-09-25
#> 4 big   2020-09-24
#> 5 huge  2020-09-27
```

```
df %>%
  select(a, !where(is.integer))
#> # A tibble: 5 x 10
#>       a group      b c      d      e      f_col      g_col      col_h      col_i
#>   <dbl> <chr> <dbl> <lgl> <chr> <fct>   <date>      <dtm>      <list>   <list>
#> 1     1  1 a      9 TRUE  A      tiny   NA      2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 2     2  4 a      NA FALSE B      small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 3     3  NA b      8 NA    C      medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
#> 4     4  3 b     10 FALSE <NA> big   2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 5     5  5 b      7 TRUE  E      huge  2020-09-27 NA      <dbl [2~ <dbl[,11] [32 x~
```

```
df %>%
  select(where(is.character) | where(is.factor))
#> # A tibble: 5 x 3
#>   group d      e
#>   <chr> <chr> <fct>
#> 1 a    A      tiny
#> 2 a    B      small
#> 3 b    C      medium
```



```
#> 4 b      <NA> big
#> 5 b      E     huge
```

```
df %>%
  select(where(~ is.double(.) | is.list()))
#> # A tibble: 5 x 6
#>       a      b f_col      g_col      col_h      col_i
#>   <dbl> <dbl> <date>    <dtm>      <list>    <list>
#> 1     1     9 NA      2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2     4     NA 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3    NA     8 2020-09-25 2020-09-26 07:12:00 <dbl [2]> <chr [5]>
#> 4     3    10 2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x 5]>
#> 5     5     7 2020-09-27 NA      <dbl [2]> <dbl[,11] [32 x 11]>
```

```
df %>%
  select_if(~ is.character(.x) | is.factor(.x))
#> # A tibble: 5 x 3
#>   group d      e
#>   <chr> <chr> <fct>
#> 1 a    A     tiny
#> 2 a    B     small
#> 3 b    C     medium
#> 4 b    <NA> big
#> 5 b    E     huge
```

4.2 filter() Rows

4.2.1 by row_number()

```
df %>%
  filter(row_number() == 1)
#> # A tibble: 1 x 10
#>   group      a      b c      d      e      f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>    <dtm>      <list>    <list>
#> 1 a      1      9 TRUE A     tiny NA      2020-09-24 02:24:00 <dbl [2]> <NULL>
```

```
df %>%
  filter(row_number() > 1)
#> # A tibble: 4 x 10
#>   group      a      b c      d      e      f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>    <dtm>      <list>    <list>
#> 1 a      4      NA FALSE B     small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 2 b      NA      8 NA    C     medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
```

```
#> 3 b          3    10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 4 b          5      7 TRUE  E    huge    2020-09-27 NA                      <dbl [2~ <dbl[,11] [32 x~
```

4.2.2 by Name

```
df %>%
  filter(a == 2)
#> # A tibble: 0 x 10
#> # ... with 10 variables: group <chr>, a <dbl>, b <dbl>, c <lgl>, d <chr>, e <fct>,
#> #   f_col <date>, g_col <dtm>, col_h <list>, col_i <list>
```

```
df %>%
  filter(a != 2)
#> # A tibble: 4 x 10
#>   group      a      b c      d      e    f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>      <dtm>      <list>   <list>
#> 1 a          1      9 TRUE  A    tiny  NA        2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 a          4      NA FALSE B    small 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 3 b          3     10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x ~
#> 4 b          5      7 TRUE  E    huge    2020-09-27 NA                      <dbl [2]> <dbl[,11] [32 x~
```

```
df %>%
  filter(c)
#> # A tibble: 2 x 10
#>   group      a      b c      d      e    f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>      <dtm>      <list>   <list>
#> 1 a          1      9 TRUE  A    tiny  NA        2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 b          5      7 TRUE  E    huge    2020-09-27 NA                      <dbl [2]> <dbl[,11] [32 x~
```

```
df %>%
  filter(!c)
#> # A tibble: 2 x 10
#>   group      a      b c      d      e    f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>      <dtm>      <list>   <list>
#> 1 a          4      NA FALSE B    small 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 2 b          3     10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2]> <df[,5] [150 x ~
```

```
df %>%
  filter(a == 5, d == "E")
#> # A tibble: 1 x 10
#>   group      a      b c      d      e    f_col      g_col      col_h      col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>      <dtm>      <list>   <list>
#> 1 b          5      7 TRUE  E    huge    2020-09-27 NA                      <dbl [2]> <dbl[,11] [32 x~
```

```
df %>%
  filter(a >= 3 | f_col == "2020-09-24")
#> # A tibble: 3 x 10
#>   group   a     b c     d     e   f_col   g_col   col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>   <list> <list>
#> 1 a         4    NA FALSE B     small 2020-09-26 2020-09-25 04:48:00 <dbl [2]> <dbl [1]>
#> 2 b         3    10 FALSE <NA> big   2020-09-24 2020-09-27 09:36:00 <dbl [2]> <dbl [5] [150 x ~
#> 3 b         5     7 TRUE  E     huge 2020-09-27 NA                <dbl [2]> <dbl [11] [32 x ~
```

```
df %>%
  filter(a < 2 | c)
#> # A tibble: 2 x 10
#>   group   a     b c     d     e   f_col   g_col   col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>   <list> <list>
#> 1 a         1     9 TRUE  A     tiny NA       2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 b         5     7 TRUE  E     huge 2020-09-27 NA                <dbl [2]> <dbl [11] [32 x ~
```

```
df %>%
  filter(!is.na(a), !is.na(b), !is.na(d))
#> # A tibble: 2 x 10
#>   group   a     b c     d     e   f_col   g_col   col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>   <list> <list>
#> 1 a         1     9 TRUE  A     tiny NA       2020-09-24 02:24:00 <dbl [2]> <NULL>
#> 2 b         5     7 TRUE  E     huge 2020-09-27 NA                <dbl [2]> <dbl [11] [32 x ~
```

4.2.3 by Type

```
df %>%
  filter(across(where(is.numeric), ~ .x >= 5))
#> # A tibble: 1 x 10
#>   group   a     b c     d     e   f_col   g_col   col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>   <list> <list>
#> 1 b         5     7 TRUE  E     huge 2020-09-27 NA                <dbl [2]> <dbl [11] [32 x ~
```

```
df %>%
  filter_if(is.numeric, ~ .x >= 5)
#> # A tibble: 1 x 10
#>   group   a     b c     d     e   f_col   g_col   col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>   <list> <list>
#> 1 b         5     7 TRUE  E     huge 2020-09-27 NA                <dbl [2]> <dbl [11] [32 x ~
```

```
df %>%
  filter_if(is.list, ~ map_lgl(.x, ~ !is.null(.x)))
```

```
#> # A tibble: 4 x 10
#>   group     a     b c     d     e   f_col   g_col           col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>           <list> <list>
#> 1 a         4     NA FALSE B     small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 2 b         NA      8 NA    C     medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
#> 3 b         3     10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 4 b         5      7 TRUE  E     huge   2020-09-27 NA                <dbl [2~ <dbl[,11] [32 x~
```

4.3 arrange() Rows

```
df %>%
  arrange(a)
#> # A tibble: 5 x 10
#>   group     a     b c     d     e   f_col   g_col           col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>           <list> <list>
#> 1 a         1      9 TRUE  A     tiny   NA       2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 2 b         3     10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 3 a         4     NA FALSE B     small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 4 b         5      7 TRUE  E     huge   2020-09-27 NA                <dbl [2~ <dbl[,11] [32 x~
#> 5 b         NA      8 NA    C     medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
```

```
df %>%
  arrange(desc(a))
#> # A tibble: 5 x 10
#>   group     a     b c     d     e   f_col   g_col           col_h   col_i
#>   <chr> <dbl> <dbl> <lgl> <chr> <fct> <date>   <dtm>           <list> <list>
#> 1 b         5      7 TRUE  E     huge   2020-09-27 NA                <dbl [2~ <dbl[,11] [32 x~
#> 2 a         4     NA FALSE B     small 2020-09-26 2020-09-25 04:48:00 <dbl [2~ <dbl [1]>
#> 3 b         3     10 FALSE <NA> big    2020-09-24 2020-09-27 09:36:00 <dbl [2~ <df[,5] [150 x ~
#> 4 a         1      9 TRUE  A     tiny   NA       2020-09-24 02:24:00 <dbl [2~ <NULL>
#> 5 b         NA      8 NA    C     medium 2020-09-25 2020-09-26 07:12:00 <dbl [2~ <chr [5]>
```